

# Marvis Osazuwa CV

## **Analytics Engineer • Data Scientist**

Manchester, UK • +44 7349 949871 • [marvis.osazuwa@hotmail.com](mailto:marvis.osazuwa@hotmail.com) • [linkedin.com/in/marvisosazuwa](https://linkedin.com/in/marvisosazuwa) • [github.com/marz1307](https://github.com/marz1307) • [marz1307.github.io](https://marz1307.github.io)

## **Profile**

Analytics Engineer and Data Scientist with production experience in B2B SaaS customer analytics. Specialist in full-lifecycle data systems: orchestrated pipelines, dimensional modelling, and applied ML with built-in explainability, in Python, SQL, and dbt. Shipped a greenfield customer intelligence data layer and independently delivered an end-to-end churn prediction stack, both in live production. MSc Data Science graduate with Distinction. Full UK work rights (Graduate visa, no sponsorship needed). EU Blue Card eligible.

## **Experience**

### **Analytics Engineer • Force24 • Leeds, UK (Hybrid) • Fixed-term contract, January 2026 – April 2026**

Marketing automation SaaS. A 16-week collaborative build of a 3-tier customer intelligence platform. Within the team delivery I owned the data layer outright as sole architect and author, was primary author across 8 backend domains, and shipped the customer-facing frontend.

- Sole architect and author of a dbt and Dagster data pipeline that consolidated 4 disconnected business systems into one analytics-ready warehouse on a deduplicated account spine (over 40 dbt models, 5 daily schedules, ~93% solo commit share). Established a single canonical customer ID across the business, which closed the record-matching gaps that had made Customer Success books unjoinable.
- Cut daily pipeline compute by ~95% by re-architecting high-volume revenue models as incremental on an append-only raw layer with deterministic hash IDs. Productionised the platform on Oracle Cloud with idempotent bootstrap, TLS, VPN allowlist, and three-level operator runbooks, so it runs 24/7 at zero infrastructure cost.
- Made support attribution auditor-defensible with a 10-step deterministic classification chain that maps every ticket to exactly one account. This replaced a black-box best-guess join and closed the last open finding in the Customer Success metrics review.
- Gated the build behind 123 dbt tests and 82 pytest tests in GitHub Actions CI, which caught 22 data-quality bugs during the first production run. Silent metric drift now fails a test instead of surfacing in the boardroom.
- Shipped the customer-facing backend and Angular UI (customer list with 10 server-side filters, dashboard and reporting endpoints, sole-author Action Tracker across 7 outcome types), which gave Customer Success a single workflow for retention actions and produced the first labelled outcomes dataset for a planned causal uplift model.

**Stack:** Dagster • dbt • PostgreSQL • Python • FastAPI • Redis • Angular • Docker • Caddy • Oracle Cloud • GitHub Actions • Git

### **CRM Data Specialist • Natural Clinic • Istanbul, Türkiye • June 2024 – January 2025**

Healthcare marketing organisation. CRM data and analytics ownership across patient acquisition, conversion, and retention.

- Rebuilt the CRM data layer end to end as the single source of truth for sales, marketing, and clinical operations. Lead response time dropped 30%; conversion rose 15%; duplicates and data errors fell 40%; email engagement rose 25%.
- Replaced weekly manual spreadsheet reporting with real-time Power BI dashboards and exec-ready reports, cutting leadership reporting latency from 7 days to live and freeing ~6 analyst-hours per

week on report assembly.

- Integrated the CRM with the marketing automation stack (email, SMS, paid funnels) and engineered customer data flows between CRM, patient management, and marketing tooling, which removed three days a week of manual reconciliation.
- Translated ambiguous stakeholder requests from marketing, clinical operations, and customer service into 12+ shipped data models and dashboards, including the lead-scoring and conversion-funnel models that drove the 15% conversion lift.

**Stack:** SQL · Python · Power BI · Zoho CRM · Marketing automation · Workflow automation

## **Strategic Data Insights Analyst · Federal Mortgage Bank of Nigeria · Abuja, Nigeria · December 2018 – November 2019**

Federal mortgage institution managing Nigeria's national housing fund. Loan portfolio analytics and IT-governance reporting in the Asset Creation Unit of the Loan Administration group.

- Contributed to a data-integrity overhaul that improved core dataset accuracy by 40% in six months, materially reducing the risk of faulty loan approvals.
- Built Python and SQL analytics on top of Tableau to track loan repayment behaviour, supporting a 15% lift in repayment-monitoring efficiency and a 25% reduction in manual data entry.
- Cut IT-asset-governance approval turnaround by 30% by re-mapping the approval workflow and instrumenting stage-level reporting in Tableau, shrinking median time-to-approve across the Loan Administration group.

**Stack:** Python · SQL · Tableau · Excel

## **Education**

**MSc Data Science** · University of Salford · Manchester, UK · January 2025 – May 2026 · Distinction Dissertation: Agentic ELT Data Platform for Customer Intelligence.

**MSc Big Data Analytics and Management** · Bahcesehir University · Istanbul, Türkiye · September 2020 – March 2023 · GPA 3.67 / 4.00 Thesis: distributed PySpark pipeline for football scouting.

**BEng Computer Engineering** · University of Benin · Edo, Nigeria · September 2011 – August 2017

## **Projects**

**Agentic ELT Data Platform for Customer Intelligence** · MSc Dissertation (Salford) · 2026 · Live B2B SaaS Environment Solo research build, end to end: I designed and wrote every layer myself, from ingestion to the serving API, against a live B2B SaaS environment under NDA. It ran on the same environment as my Force24 contract but is independent work; the Force24 platform itself was a team delivery.

- Ingestion and warehouse: custom Python extractors ingested over 1M records from multiple vendor APIs into a JSONB-first PostgreSQL warehouse, orchestrated with Dagster.
- Serving and agent access: FastAPI service with JWT auth and Row-Level Security, Angular dashboard, and a Model Context Protocol (MCP) endpoint queryable by any MCP-compatible LLM agent for per-account risk profiling, driver explanations, and intervention recommendations. Output prioritisation surfaced high-risk accounts to CSMs for retention action.
- Modelling: 48 dbt models across raw, staging, intermediate, and marts; layered Kimball-style dimensional design.
- Intelligence stack: survival analysis (cross-validated C-index ~0.94 on held-out cohort) for time-to-churn; gradient-boosted classification with SHAP explanations (AUC ~0.95 vs ~0.89 logistic baseline) for churn likelihood and drivers; DR-Learner causal inference (econml) correcting reactive-assignment selection bias by ~50 percentage points on the naïve estimator for honest CSM treatment effect.

Python PostgreSQL JSONB dbt Dagster PostgresML XGBoost SHAP Causal Inference Survival Analysis FastAPI Angular MCP

**Pharmaceutical Side Effect Classification** · Side Project · 2025 · Healthcare ML Production-shape multi-class classifier mapping free-text adverse-event descriptions to ten clinical categories across 11,825 marketed medicines. Random Forest at 98.5% accuracy on held-out test, Logistic Regression baseline at 97.2% (interpretable comparison within two points). Above 95% F1 across every retained category. Inference CLI and pytest test suite shipped. Python scikit-learn TF-IDF Random Forest Logistic Regression pytest

**Big Data Player Scouting (MSc Thesis)** · Bahcesehir Thesis · 2023 · Sports / Big Data Distributed PySpark pipeline ranking and recommending footballers across five top European leagues using UEFA event data

and PlayeRank methodology. 500+ players, hypothesis-driven research. Published thesis, public codebase.  
Python PySpark Big Data Recommender Systems

**Equity Forecasting** · Side Project · 2025 · Financial Time Series End-to-end R analysis package for univariate equity forecasting on 5,124 daily NYSE observations. ARIMA, seasonal ARIMA, and ETS through a MODEL\_REGISTRY pattern, with formal residual diagnostics (Ljung-Box, Shapiro-Wilk) and stationarity tests (ADF, KPSS) combined into a single decision rule. R ARIMA ETS Time Series Statistical Testing

## Technical Skills

**Languages:** SQL · Python (pandas, NumPy, scikit-learn) · R · PySpark · Bash **Warehouse & pipelines:** PostgreSQL (JSONB + GIN) · SQL Server / T-SQL · dbt · Dagster · Kimball dimensional modelling · ELT/ETL design **Distributed & lakehouse:** PySpark · Spark MLlib · MLflow · Databricks **ML & statistics:** XGBoost · SHAP · scikit-survival · econml (DR-Learner) · PostgresML · causal inference · survival analysis · RoBERTa and transformer sentiment (NLP) · TF-IDF · LDA · ARIMA/ETS time series **Serving & agentic access:** FastAPI · REST APIs · Model Context Protocol (MCP) endpoints for LLM agents · Docker · Caddy · Oracle Cloud **BI:** Power BI · Tableau **CI & source control:** Git · GitHub Actions · pytest

**Working knowledge (coursework and certifications):** Snowflake · BigQuery (Google Cloud certified) · Airflow · Looker · AWS · GCP · Azure

## Languages

**English:** Native · **German:** B1

## Certifications

Engineer Data for Predictive Modeling with BigQuery ML (Google Cloud) · Python for Data Science, AI & Development (Coursera) · Generative AI for Business Leaders (Coursera) · Enterprise Design Thinking Practitioner / Team Essentials for AI / Co-Creator (IBM)

## Community & Leadership

**Microsoft Learn Student Ambassador** · Bahçeşehir University · Istanbul, Türkiye · January 2021 – September 2023 (2 yrs 9 mo) Cross-university Azure workshops and a community hackathon with fellow ambassadors and interested learners.