

Marvis Osazuwa

Analytics Engineer · Data Scientist

Manchester, UK · +44 7349 949871 · marvis.osazuwa@hotmail.com · linkedin.com/in/marvisosazuwa · github.com/marz1307 · marvisosazuwa.com

PROFILE

Analytics Engineer and Data Scientist with production experience in B2B SaaS customer analytics. Specialist in full-lifecycle data systems: orchestrated pipelines, dimensional modelling, and applied ML with built-in explainability, in Python, SQL, and dbt. Shipped a greenfield customer intelligence data layer and independently delivered an end-to-end churn prediction stack, both in live production. MSc Data Science graduate with Distinction. Full UK work rights (Graduate visa, no sponsorship needed). EU Blue Card eligible.

EXPERIENCE

Analytics Engineer · Force24 · Leeds, UK (Hybrid)

Fixed-term contract, January 2026 to April 2026

A defined 16-week build in which I was sole architect and author of the data layer for a 3-tier customer intelligence platform, primary author across 8 backend domains, and shipping engineer for the customer-facing frontend.

- Sole architect and author of a dbt and Dagster data pipeline that consolidated 4 disconnected business systems into one analytics-ready warehouse on a deduplicated account spine: over 40 dbt models, 5 daily schedules, ~93% solo commit share. Established a single canonical customer ID across the business, eliminating record-matching gaps that had previously made Customer Success books unjoinable.
- Cut daily pipeline compute by ~95% by re-architecting high-volume revenue models as incremental on an append-only raw layer with deterministic hash IDs. Productionised the platform on Oracle Cloud with idempotent bootstrap, TLS, VPN allowlist, and three-level operator runbooks: 24/7 platform at zero infrastructure cost.
- Lifted support-attribution accuracy to auditor-defensible by designing a 10-step deterministic classification chain mapping every ticket to a single account, replacing a black-box best-guess join and closing the last open finding in the Customer Success metrics review.
- Locked correctness end-to-end with 123 dbt tests and 82 pytest tests gated through GitHub Actions CI, surfacing and resolving 22 data-quality bugs during the first production build. Silent metric drift now caught in CI instead of the boardroom.
- Shipped the customer-facing backend and Angular UI (customer list with 10 server-side filters, dashboard and reporting endpoints, sole-author Action Tracker across 7 outcome types) giving Customer Success a single workflow for retention actions and generating the first labelled outcomes dataset for a planned causal uplift model. Designed on the principle that dashboards should drive action, not just display data.

Stack: Dagster, dbt, PostgreSQL, Python, FastAPI, Redis, Angular, Docker, Caddy, Oracle Cloud, GitHub Actions, Git.

CRM Data Specialist · Natural Clinic · Istanbul, Türkiye

June 2024 – January 2025

Healthcare marketing organisation. CRM data and analytics ownership across patient acquisition, conversion, and retention.

- Rebuilt the CRM data layer end to end as the single source of truth for sales, marketing, and clinical operations. Lead response time dropped 30%; conversion rose 15%; duplicates and data errors fell 40%; email engagement rose 25%.
- Replaced weekly manual spreadsheet reporting with real-time Power BI dashboards and exec-ready reports, cutting leadership reporting latency from 7 days to live and freeing ~6 analyst-hours per week on report assembly.
- Integrated the CRM with the marketing automation stack (email, SMS, paid funnels) and engineered customer data flows between CRM, patient management, and marketing tooling, eliminating three days a week of manual reconciliation.
- Translated ambiguous stakeholder requests from marketing, clinical operations, and customer service into 12+ shipped data models and dashboards, including the lead-scoring and conversion-funnel models that drove the 15% conversion lift.

Stack: SQL, Python, Power BI, Zoho CRM, marketing automation, workflow automation.

Strategic Data Insights Analyst · Federal Mortgage Bank of Nigeria · Abuja, Nigeria

December 2018 – November 2019

Federal mortgage institution managing Nigeria's national housing fund. Loan portfolio analytics and IT-governance reporting in the Asset Creation Unit of the Loan Administration group, under senior supervision.

- Contributed to a data-integrity overhaul that improved core dataset accuracy by 40% in six months, materially reducing the risk of faulty loan approvals.
- Built Python and SQL analytics on top of Tableau to track loan repayment behaviour, supporting a 15% lift in repayment-monitoring efficiency and a 25% reduction in manual data entry.
- Cut IT-asset-governance approval turnaround by 30% by re-mapping the approval workflow and instrumenting stage-level reporting in Tableau, shrinking median time-to-approve across the Loan Administration group.

Stack: Python, SQL, Tableau, Excel.

EDUCATION

MSc Data Science · University of Salford · Manchester, UK

January 2025 – May 2026 · Distinction · Dissertation: Agentic ELT Data Platform for Customer Intelligence

MSc Big Data Analytics and Management · Bahcesehir University · Istanbul, Türkiye

September 2020 – March 2023 · GPA 3.67 / 4.00 · Thesis: distributed PySpark pipeline for football scouting

BEng Computer Engineering · University of Benin · Edo, Nigeria

September 2011 – August 2017

PROJECTS

Agentic ELT Data Platform for Customer Intelligence

MSc Dissertation (Salford) · 2026 · Live B2B SaaS Environment

End-to-end ELT platform and three-model churn intelligence stack for a live B2B SaaS environment under NDA. Built a custom Python ingestion layer that landed over 1M records from multiple vendor APIs into a JSONB-first PostgreSQL warehouse via Dagster, modelled the

data through 48 dbt models, and shipped a three-model intelligence stack: survival analysis (when accounts churn), gradient-boosted classification with SHAP explanations (how likely they are to churn, with drivers), and causal inference correcting reactive-assignment selection bias in CSM effectiveness estimates. Surfaced via a FastAPI service with JWT auth and Row-Level Security, an Angular dashboard, and a Model Context Protocol (MCP) endpoint exposing the entire intelligence layer to agentic LLM workflows. Identified high-risk accounts for prioritised CSM retention action.

Python PostgreSQL JSONB dbt Dagster PostgresML XGBoost SHAP Causal Inference Survival Analysis FastAPI Angular MCP

Pharmaceutical Side Effect Classification

Side Project · 2025 · Healthcare ML

Production-grade multi-class classifier mapping free-text adverse-event descriptions to ten clinical categories across 11,825 marketed medicines. Random Forest at 98.5 percent accuracy on held-out test, Logistic Regression baseline at 97.2 percent (interpretable comparison within two points). Above 95 percent F1 across every retained category. Inference CLI and pytest test suite shipped.

Python scikit-learn TF-IDF Random Forest Logistic Regression pytest

Big Data Player Scouting (MSc Thesis)

Bahcesehir Thesis · 2023 · Sports / Big Data

Distributed PySpark pipeline ranking and recommending footballers across five top European leagues using UEFA event data and PlayeRank methodology. 500 plus players, hypothesis-driven research. Published thesis, public codebase.

Python PySpark Big Data Recommender Systems

Equity Forecasting

Side Project · 2025 · Financial Time Series

End-to-end R analysis package for univariate equity forecasting on 5,124 daily NYSE observations. ARIMA, seasonal ARIMA, and ETS through a MODEL_REGISTRY pattern, with formal residual diagnostics (Ljung-Box, Shapiro-Wilk) and stationarity tests (ADF, KPSS) combined into a single decision rule.

R ARIMA ETS Time Series Statistical Testing

TECHNICAL SKILLS

Languages: SQL, Python, R, PySpark, T-SQL, Bash
Transformation & Orchestration: dbt, Dagster, Apache Spark, Airflow
Warehousing & Storage: PostgreSQL (JSONB + GIN), Snowflake, BigQuery, Databricks, SQL Server
Service Layer: FastAPI, REST APIs, Docker, Caddy, Model Context Protocol (MCP)
Source Control & CI: Git, GitHub, GitHub Actions, pytest
Cloud: AWS, GCP, Azure, Oracle Cloud
ML & Statistics: scikit-learn, XGBoost, PostgresML, scikit-survival, econml (DR-Learner), SHAP, MLflow, RoBERTa, ARIMA, Time Series Analysis, NLP, Sentiment Analysis, TF-IDF, LDA, Causal Inference, Survival Analysis, A/B testing / experimentation
Agentic AI: Model Context Protocol (MCP), agentic ELT, LLM-accessible data layers
BI & Visualisation: Power BI, Tableau, Looker

LANGUAGES

- **English:** Native
- **German:** B1

CERTIFICATIONS

- Engineer Data for Predictive Modeling with BigQuery ML · Google Cloud
- Python for Data Science, AI & Development · Coursera
- Generative AI for Business Leaders · Coursera
- Enterprise Design Thinking Practitioner / Team Essentials for AI / Co-Creator · IBM

COMMUNITY & LEADERSHIP

Microsoft Learn Student Ambassador · Bahçeşehir University · Istanbul, Türkiye

September 2020 – June 2023 (2 yrs 9 mo) · Ran Azure workshops, organised a campus hackathon, grew the Bahçeşehir Microsoft Learn chapter.

Available immediately (MSc completed May 2026). Targeting Analytics Engineer / Data Scientist roles in the UK and Germany/EU. UK Graduate visa: no sponsorship needed; EU Blue Card eligible.