Marvis Osazuwa

Analytics Engineer · Data Platform

Location: Manchester, UK · **Phone:** +44 7349 949871 · **Email:** marvis.osazuwa@hotmail.com · **LinkedIn:** linkedin.com/in/marvisosazuwa · **GitHub:** github.com/marz1307 · **Portfolio:** marvisosazuwa.com

PROFILE

Analytics Engineer who builds the warehouse as a product: sole architect of Force24's customer intelligence data layer — 40+ dbt models in Kimball design on a deduplicated account spine, 123 dbt and 82 pytest tests that caught 22 data-quality bugs in CI, daily pipeline compute cut ~95%, productionised on Oracle Cloud at zero infrastructure cost — with applied ML layered on top where it pays off (XGBoost with SHAP, survival analysis). The through-line: I make data trustworthy, then I make it predictive — marts gated by tests in CI, metrics with documented lineage, and the warehouse itself made queryable by LLM agents through a Model Context Protocol (MCP) endpoint. Currently building independent enterprise-grade data projects on the same stack. UK Graduate visa: full right to work in the UK, no sponsorship needed; EU Blue Card eligible for Germany and the EU.

EXPERIENCE

Analytics Engineer · Force24 · Leeds, UK (Hybrid)

Fixed-term contract · January 2026 – April 2026

Marketing automation SaaS. A 16-week collaborative build of a 3-tier customer intelligence platform. Within the team delivery I owned the data layer outright as sole architect and author, was primary author across 8 backend domains, and shipped the customer-facing frontend.

- Sole architect and author of a dbt and Dagster data pipeline that consolidated 4 disconnected business systems into one analytics-ready warehouse on a deduplicated account spine (over 40 dbt models, 5 daily schedules, ~93% solo commit share). Established a single canonical customer ID across the business, which closed the record-matching gaps that had made Customer Success books unjoinable.
- Cut daily pipeline compute by ~95% by re-architecting high-volume revenue models as incremental on an append-only raw layer with deterministic hash IDs. Productionised the platform on Oracle Cloud with idempotent bootstrap, TLS, VPN allowlist, and three-level operator runbooks, so it runs 24/7 at zero infrastructure cost.
- Made support attribution auditor-defensible with a 10-step deterministic classification chain that maps every ticket to exactly one account. This replaced a black-box best-guess join and closed the last open finding in the Customer Success metrics review.
- Gated the build behind 123 dbt tests and 82 pytest tests in GitHub Actions CI, which caught 22 data-quality bugs during the first production run. Silent metric drift now fails a test instead of surfacing in the boardroom.
- Shipped the customer-facing backend and Angular UI (customer list with 10 server-side filters, dashboard and reporting endpoints, sole-author Action Tracker across 7 outcome types), which gave Customer Success a single workflow for retention actions and produced the first labelled outcomes dataset for a planned causal uplift model.

Stack: `Dagster, dbt, PostgreSQL, Python, FastAPI, Redis, Angular, Docker, Caddy, Oracle Cloud, GitHub Actions, Git.`

CRM Data Specialist · Natural Clinic · Istanbul, Türkiye

June 2024 – January 2025

Healthcare marketing organisation. CRM data and analytics ownership across patient acquisition, conversion, and retention.

- Rebuilt the CRM data layer end to end as the single source of truth for sales, marketing, and clinical operations. Lead response time dropped 30%; conversion rose 15%; duplicates and data errors fell 40%; email engagement rose 25%.
- Replaced weekly manual spreadsheet reporting with real-time Power BI dashboards and exec-ready reports, cutting leadership reporting latency from 7 days to live and freeing ~6 analyst-hours per week on report assembly.
- Integrated the CRM with the marketing automation stack (email, SMS, paid funnels) and engineered customer data flows between CRM, patient management, and marketing tooling, which removed three days a week of manual reconciliation.
- Translated ambiguous stakeholder requests from marketing, clinical operations, and customer service into 12+ shipped data models and dashboards, including the lead-scoring and conversion-funnel models that drove the 15% conversion lift.

Stack: `SQL, Python, Power BI, Zoho CRM, marketing automation, workflow automation.`

Strategic Data Insights Analyst · Federal Mortgage Bank of Nigeria · Abuja, Nigeria

December 2018 – November 2019

Federal mortgage institution managing Nigeria's national housing fund. Loan portfolio analytics and IT-governance reporting in the Asset Creation Unit of the Loan Administration group.

- Contributed to a data-integrity overhaul that improved core dataset accuracy by 40% in six months, materially reducing the risk of faulty loan approvals.
- Built Python and SQL analytics on top of Tableau to track loan repayment behaviour, supporting a 15% lift in repayment-monitoring efficiency and a 25% reduction in manual data entry.
- Cut IT-asset-governance approval turnaround by 30% by re-mapping the approval workflow and instrumenting stage-level reporting in Tableau, shrinking median time-to-approve across the Loan Administration group.

Stack: `Python, SQL, Tableau, Excel.`

EDUCATION

MSc Data Science · University of Salford · Manchester, UK

January 2025 – May 2026 · Distinction · Dissertation: Agentic ELT Data Platform for Customer Intelligence

MSc Big Data Analytics and Management · Bahcesehir University · Istanbul, Türkiye

September 2020 – March 2023 · GPA 3.67 / 4.00 · Thesis: distributed PySpark pipeline for football scouting

BEng Computer Engineering · University of Benin · Edo, Nigeria

September 2011 – August 2017

PROJECTS

Agentic ELT Data Platform for Customer Intelligence

MSc Dissertation (Salford) · 2026 · Live B2B SaaS Environment

Solo research build, end to end: every layer designed and written by me alone, from ingestion to the serving API, against a live B2B SaaS environment under NDA. Ran on the same environment as the Force24 contract but is independent work; the Force24 platform itself was a team delivery.

- **Ingestion and warehouse:** custom Python extractors ingested over 1M records from multiple vendor APIs into a JSONB-first PostgreSQL warehouse, orchestrated with Dagster.

- **Serving and agent access:** FastAPI service with JWT auth and Row-Level Security, Angular dashboard, and a Model Context Protocol (MCP) endpoint queryable by any MCP-compatible LLM agent for per-account risk profiling, driver explanations, and intervention recommendations. Output prioritisation surfaced high-risk accounts to CSMs for retention action.
- **Modelling:** 48 dbt models across raw, staging, intermediate, and marts; layered Kimball-style dimensional design.
- **Intelligence stack:** survival analysis (cross-validated C-index ~0.94 on held-out cohort) for time-to-churn; gradient-boosted classification with SHAP explanations (AUC ~0.95 vs ~0.89 logistic baseline) for churn likelihood and drivers; DR-Learner causal inference (econml) correcting reactive-assignment selection bias by ~50 percentage points on the naïve estimator for honest CSM treatment effect.

`Python` `PostgreSQL` `JSONB` `dbt` `Dagster` `PostgresML` `XGBoost` `SHAP` `Causal Inference` `Survival Analysis` `FastAPI` `Angular` `MCP`

Equity Forecasting

Side Project · 2025 · Financial Time Series

End-to-end R analysis package for univariate equity forecasting on 5,124 daily NYSE observations. ARIMA, seasonal ARIMA, and ETS through a MODEL_REGISTRY pattern, with formal residual diagnostics (Ljung-Box, Shapiro-Wilk) and stationarity tests (ADF, KPSS) combined into a single decision rule.

`R` `ARIMA` `ETS` `Time Series` `Statistical Testing`

Pharmaceutical Side Effect Classification

Side Project · 2025 · Healthcare ML

Production-shape multi-class classifier mapping free-text adverse-event descriptions to ten clinical categories across 11,825 marketed medicines. Random Forest at 98.5% accuracy on held-out test, Logistic Regression baseline at 97.2% (interpretable comparison within two points). Above 95% F1 across every retained category. Inference CLI and pytest test suite shipped.

`Python` `scikit-learn` `TF-IDF` `Random Forest` `Logistic Regression` `pytest`

Big Data Player Scouting (MSc Thesis)

Bahcesehir Thesis · 2023 · Sports / Big Data

Distributed PySpark pipeline ranking and recommending footballers across five top European leagues using UEFA event data and PlayeRank methodology. 500+ players, hypothesis-driven research. Published thesis, public codebase.

`Python` `PySpark` `Big Data` `Recommender Systems`

TECHNICAL SKILLS

Languages: SQL, Python (pandas, NumPy, scikit-learn), R, PySpark, T-SQL, Bash
Transformation & Orchestration: dbt, Dagster, Apache Spark, Airflow
Warehousing & Storage: PostgreSQL (JSONB + GIN), Snowflake, BigQuery, Databricks, SQL Server
Service Layer: FastAPI, REST APIs, Docker, Caddy, Model Context Protocol (MCP)
Source Control & CI: Git, GitHub, GitHub Actions, pytest
Cloud: AWS, GCP, Azure, Oracle Cloud
ML & Statistics: scikit-learn, XGBoost, PostgresML, scikit-survival, econml (DR-Learner), SHAP, MLflow, RoBERTa, ARIMA, Time Series Analysis, NLP, Sentiment Analysis, TF-IDF, LDA, Causal Inference, Survival Analysis, A/B testing / experimentation
Agentic AI: Model Context Protocol (MCP), agentic ELT, LLM-accessible data layers
BI & Visualisation: Power BI, Tableau, Looker

LANGUAGES

- **English:** Native
- **German:** B1

CERTIFICATIONS

- Engineer Data for Predictive Modeling with BigQuery ML · Google Cloud
- Python for Data Science, AI & Development · Coursera
- Generative AI for Business Leaders · Coursera
- Enterprise Design Thinking Practitioner / Team Essentials for AI / Co-Creator · IBM

COMMUNITY & LEADERSHIP

Microsoft Learn Student Ambassador · Bahçeşehir University · Istanbul, Türkiye

January 2021 – September 2023 (2 yrs 9 mo) · Cross-university Azure workshops and a community hackathon with fellow ambassadors and interested learners.

Availability: available immediately for UK roles (UK Graduate visa, no sponsorship needed) and for Germany and the EU (EU Blue Card eligible).